

## Parallel Programming HW2 Problems

- 25) L2 says if method call  $m$  precedes method call  $n$  in  $H$ , then the same is true in  $S$ .  $H$  is a history &  $S$  is a legal sequential history that is equivalent to the complete extension of  $H$ . L2 is the property that helps linearizability be compositional, so by taking away the L2 statement from the linearizability definition it takes away the compositional property. Thus, it makes the resulting property the same as sequential consistency.
- 29) The statement is NOT equivalent. A wait free object guarantees that every call finishes in a finite number of steps. In an infinite history where each thread completes an infinite number of method calls, there is no guarantee of the previous statement. Therefore, it is not equivalent.
- 30) The statement is equivalent. Lock-free objects guarantee that infinitely often, some method call finishes in a finite number of steps. In an infinite history of  $x$ , where an infinite # of method calls are completed there is always some method call being finished. therefore since it finished then it completes in a finite # of steps and there are an infinite number of calls which have this characteristic. Thus, infinitely often some method call finishes in a finite number of steps so object  $x$  is lock free.



31) A wait-free method guarantees that every call completes in a finite number of steps. A bounded wait-free method puts a bound on the # of steps a call can take. Here, the method has a bounded wait-free property because the call returns after  $2^i$  steps. This means that there are a finite # of steps the method completes, and thus the method is wait-free. Also, the call does not exceed  $2^i$  steps so therefore it has a bound.

$\therefore$  It is a bounded wait-free method

32) Thread 1 enqueues  $x$  onto the queue, & concurrently thread 2 enqueues  $y$ . Then, thread 3 dequeues

- 1 executes `tail.getAndIncrement` which sets  $i$  to 0 & the tail goes to one.
  - 2 also executes `tail.getAndIncrement` which sets  $i$  to 1 & tail goes to two.
  - 2 stores  $y$  at index 1 of items
  - 3 dequeues & finds the item at index 0 to be null  
↳ it goes to index 1 & finds & returns  $y$
  - 1 stores  $x$  at 0<sup>th</sup> index
  - 3 dequeues & finds item at 0 to be  $x$  & returns
- $\therefore$  the code returns  $y$  before  $x$  even though thread 1 executed line 15 before thread 2. So it can't be a linearization point



- line 16 example:

- thread 1 executes `tail.getAndIncrement` before 2
  - 2 then executes `item[i].set(y)` before 1 set `item[i]` to x
  - 3 dequeues 0<sup>th</sup> index & returns x
  - 3 dequeues 1<sup>st</sup> index & returns y
- ∴ line 16 is not a linearization point

- this doesn't mean enqueue is not linearizable.  
It means sometimes there is not a fixed system statement which is a linearization point.

# Parallel Programming HW2 Write Up

Andrew McCarthy

Source files are located in the hw2/src directory. My code is in main.c, lamport.c/h, and atomic.h. Instructions on how to compile and run the code is under the 'Details' section

- Overview

This program is designed to mimic the functions of a firewall in both parallel and serial execution. The parallel version enqueues packets into a queue (each source gets its own queue) while worker threads concurrently dequeue and process the packets. The serial version of the code retrieves packet from each source and processes them serially using a double for loop. The program also includes a serial-queue version which enqueues and dequeues packets before they are processed serially (this is designed to test the overhead of the queue structure). The program takes in 7 arguments when run normally. These are the number of threads (including master), the number of packets from each source, the queue depth, the mean/expected work to process each packet, the seed, the packet generation type (constant, uniform, and exponential), and the running mode (parallel, serial, and serial-queue). The program can also run a series of performance tests by making the first input parameter test and the second the seed.

- Design

The project code is in the Lamport.c/h, atomic.h and main.c files. The lamport files are where the code for the lamport/lock-free-queues are. Atomic is where the code to make use of the atomic gcc builtins is. Main.c is the where the dispatcher and worker code is as well as all the code for managing input arguments and running the program in serial, serial-queue or parallel mode.

The lamport queue has three functions: create\_queue, enqueue\_queue, and dequeue\_queue. The depth of the queue is set during allocation in the create\_queue function. Element\_t is a structure used to hold the address of enqueued elements and an update count for the slot. Queue\_t contains the tail and head information, the depth information, and a list of the elements stored in it.

The enqueue function starts by taking a copy of the tail pointer and the element at said pointer. Next, the function makes sure the element's pointer is correctly pointing to the tail of the queue. Next, the function checks if the queue is full; if it is, the function returns 0. Next, the function ensures the head pointer is not lagging behind. Next, the function checks if the current slot's pointer is non-null. This ensures that the tail pointer does not lag behind. Next, the data in the slot is copied and a DWCAS operation is performed. If it is successful, the tail pointer is incremented and 1 is returned. If it fails, the operations are restarted. The dequeue function is similar to the enqueue function except it removes an element from the queue.

The atomic.h file has two functions: CAS and DWCAS (compare-and-swap and double-compare-and-swap). These functions are used by the queue to synchronize and ensure that the tail and head pointers do not lag behind.

The dispatcher for the packets works by using an array representing the total number of packets still needed for each source, and an int representing the total number of packets that need to be generated ( $\text{numQueues} * \text{packetsPerSource}$ ). The dispatcher continuously loops through each queue and generates and inserts a packet if the queue is not full. If a packet is successfully inserted, the corresponding value in the packets-needed array is decremented and the total-number-of-packets int is decremented. Once the packets-needed int is 0, the loop exits.

The worker threads continuously dequeue elements from the queue and calculate the fingerprint. Once the worker has processed the designated number of packets (number of packets that each source generates) it exits.

The serial and serial-queue versions of the code work by looping through all the sources then generating and processing the necessary packets. The serial-queue version enqueues and dequeues packets before processing them.

- Divergence From Design Document

There are two big changes I made to my design after I began coding. First, I incorrectly stated how to properly use memory fences in my design document. Once I started coding I realized my mistake and changed my design plan. Second, in my design document I stated that I would load all the packets into the queues before starting the worker threads. After reading the project description again I realized this would be impossible as the queue depth could not hold all of the packets at once. I corrected this mistake once I began coding.

- Details

To create the executables for the program run `make`. To remove the executable and object files run `make clean`. To run the designated tests, make “test” the first input parameter and a seed value the second. To run in normal mode, you must provide the following parameters preceded by their corresponding flag: -n (number of threads including master), -t (packets from each source), -d (queue depth), -w (mean work per packet), -g (packet generation mode-> must be “constant”, “uniform”, “exponential”), -m (parallel or serial execution -> must be “parallel”, “serial”, “serial-queue”). A -s (seed) parameter can also be provided but it is not required. The parameters can be given in any order. When running in normal mode, the program prints out the run time to stdout. When running in test mode, the program prints the test results to stdout. Here is an example execution: `./lamport -n 14 -t 60000 -d 32 -w 3000 -s 50 -g constant -m parallel`

- Issues

One issue I ran into was the -Wstrict-aliasing warning blocking me from compiling my code. Upon looking into this warning, I learned it wasn't an issue for my use case (converting pointers to my element\_t struct in lamport.c) so I disabled it using compiler flags.

- Correctness testing

My worker threads can only finished once they process as many packets as is set by the T parameter. Due to this, if any errors caused a packet to not be properly enqueued/dequeued, the program would spin indefinitely. To test this, I ran the parallel version of my code with various input parameters and made sure it finished.

- Performance Testing

To test performance difference between serial and parallel executions, I ran the tests outlined in the project description with preset values and printed the results to stdout. The results of these tests are listed below:

--Output:

speedup queue overhead with w = 25 and n = 2 is 2.154754  
speedup queue overhead with w = 25 and n = 9 is 1.071721  
speedup queue overhead with w = 25 and n = 14 is 1.555466  
speedup queue overhead with w = 50 and n = 2 is 1.347470  
speedup queue overhead with w = 50 and n = 9 is 1.344514  
speedup queue overhead with w = 50 and n = 14 is 1.363018  
speedup queue overhead with w = 100 and n = 2 is 1.211009  
speedup queue overhead with w = 100 and n = 9 is 1.223723  
speedup queue overhead with w = 100 and n = 14 is 1.227577  
speedup queue overhead with w = 200 and n = 2 is 1.127900  
speedup queue overhead with w = 200 and n = 9 is 1.120298  
speedup queue overhead with w = 200 and n = 14 is 1.124955  
speedup queue overhead with w = 400 and n = 2 is 1.077529  
speedup queue overhead with w = 400 and n = 9 is 1.068528  
speedup queue overhead with w = 400 and n = 14 is 1.058897  
speedup queue overhead with w = 800 and n = 2 is 1.006073  
speedup queue overhead with w = 800 and n = 9 is 1.032700  
speedup queue overhead with w = 800 and n = 14 is 1.039192

dispatcher rate time with w = 1 and n = 2 is 4559.582031

dispatcher rate time with w = 1 and n = 3 is 5008.192871

dispatcher rate time with  $w = 1$  and  $n = 5$  is 4600.687012  
dispatcher rate time with  $w = 1$  and  $n = 9$  is 4278.637207  
dispatcher rate time with  $w = 1$  and  $n = 14$  is 4129.667969  
dispatcher rate time with  $w = 1$  and  $n = 28$  is 5822.081055  
speedup constant packets with  $w = 1000$  and  $n = 2$  is 1.116477  
speedup constant packets with  $w = 1000$  and  $n = 3$  is 1.153973  
speedup constant packets with  $w = 1000$  and  $n = 5$  is 1.318144  
speedup constant packets with  $w = 1000$  and  $n = 9$  is 1.613507  
speedup constant packets with  $w = 1000$  and  $n = 14$  is 1.613503  
speedup constant packets with  $w = 1000$  and  $n = 28$  is 1.123558  
speedup constant packets with  $w = 2000$  and  $n = 2$  is 2.204270  
speedup constant packets with  $w = 2000$  and  $n = 3$  is 2.223579  
speedup constant packets with  $w = 2000$  and  $n = 5$  is 2.687395  
speedup constant packets with  $w = 2000$  and  $n = 9$  is 3.012343  
speedup constant packets with  $w = 2000$  and  $n = 14$  is 3.114397  
speedup constant packets with  $w = 2000$  and  $n = 28$  is 2.656482  
speedup constant packets with  $w = 4000$  and  $n = 2$  is 4.075189  
speedup constant packets with  $w = 4000$  and  $n = 3$  is 4.466882  
speedup constant packets with  $w = 4000$  and  $n = 5$  is 5.387151  
speedup constant packets with  $w = 4000$  and  $n = 9$  is 5.929618  
speedup constant packets with  $w = 4000$  and  $n = 14$  is 6.370973  
speedup constant packets with  $w = 4000$  and  $n = 28$  is 4.855897  
speedup constant packets with  $w = 8000$  and  $n = 2$  is 8.366705  
speedup constant packets with  $w = 8000$  and  $n = 3$  is 9.396678  
speedup constant packets with  $w = 8000$  and  $n = 5$  is 11.555341  
speedup constant packets with  $w = 8000$  and  $n = 9$  is 11.704694  
speedup constant packets with  $w = 8000$  and  $n = 14$  is 12.350133  
speedup constant packets with  $w = 8000$  and  $n = 28$  is 7.593753

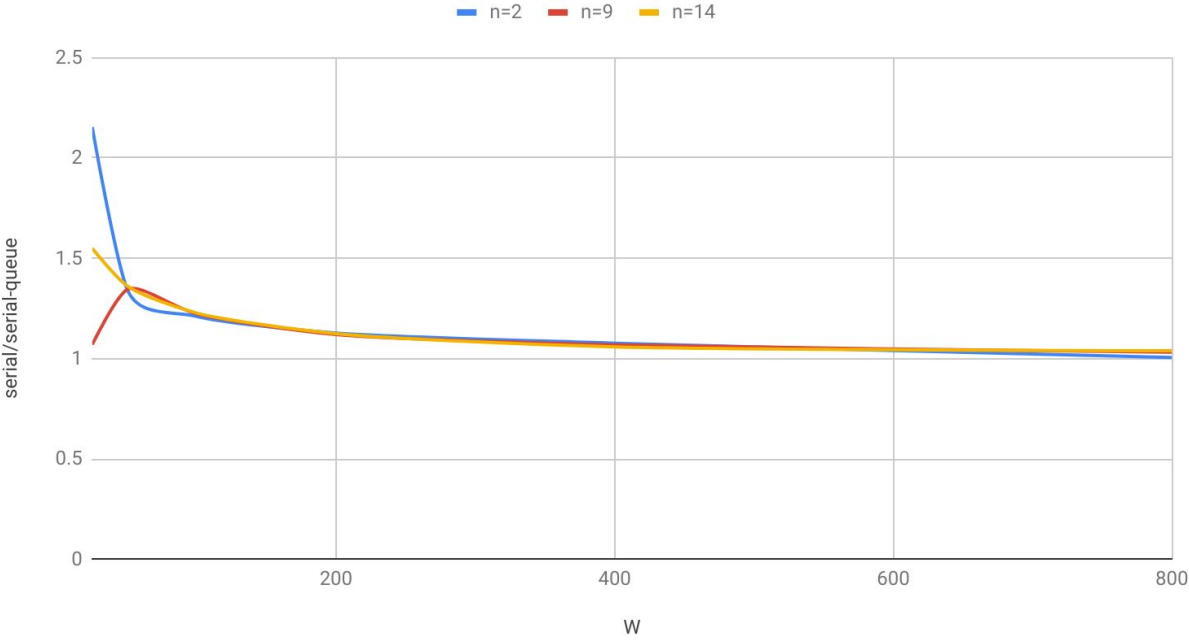
speedup uniform packets with  $w = 1000$  and  $n = 2$  is 3.316725  
speedup uniform packets with  $w = 1000$  and  $n = 3$  is 4.108906  
speedup uniform packets with  $w = 1000$  and  $n = 5$  is 4.283566  
speedup uniform packets with  $w = 1000$  and  $n = 9$  is 5.362582  
speedup uniform packets with  $w = 1000$  and  $n = 14$  is 5.106752  
speedup uniform packets with  $w = 1000$  and  $n = 28$  is 3.702444  
speedup uniform packets with  $w = 2000$  and  $n = 2$  is 6.911411  
speedup uniform packets with  $w = 2000$  and  $n = 3$  is 8.140733  
speedup uniform packets with  $w = 2000$  and  $n = 5$  is 9.215037  
speedup uniform packets with  $w = 2000$  and  $n = 9$  is 9.934791

speedup uniform packets with  $w = 2000$  and  $n = 14$  is 10.678123  
speedup uniform packets with  $w = 2000$  and  $n = 28$  is 7.932747  
speedup uniform packets with  $w = 4000$  and  $n = 2$  is 13.528705  
speedup uniform packets with  $w = 4000$  and  $n = 3$  is 15.826137  
speedup uniform packets with  $w = 4000$  and  $n = 5$  is 18.915048  
speedup uniform packets with  $w = 4000$  and  $n = 9$  is 20.077365  
speedup uniform packets with  $w = 4000$  and  $n = 14$  is 20.721141  
speedup uniform packets with  $w = 4000$  and  $n = 28$  is 13.243738  
speedup uniform packets with  $w = 8000$  and  $n = 2$  is 30.575239  
speedup uniform packets with  $w = 8000$  and  $n = 3$  is 32.185424  
speedup uniform packets with  $w = 8000$  and  $n = 5$  is 39.510968  
speedup uniform packets with  $w = 8000$  and  $n = 9$  is 40.513812  
speedup uniform packets with  $w = 8000$  and  $n = 14$  is 41.673233  
speedup uniform packets with  $w = 8000$  and  $n = 28$  is 32.894543

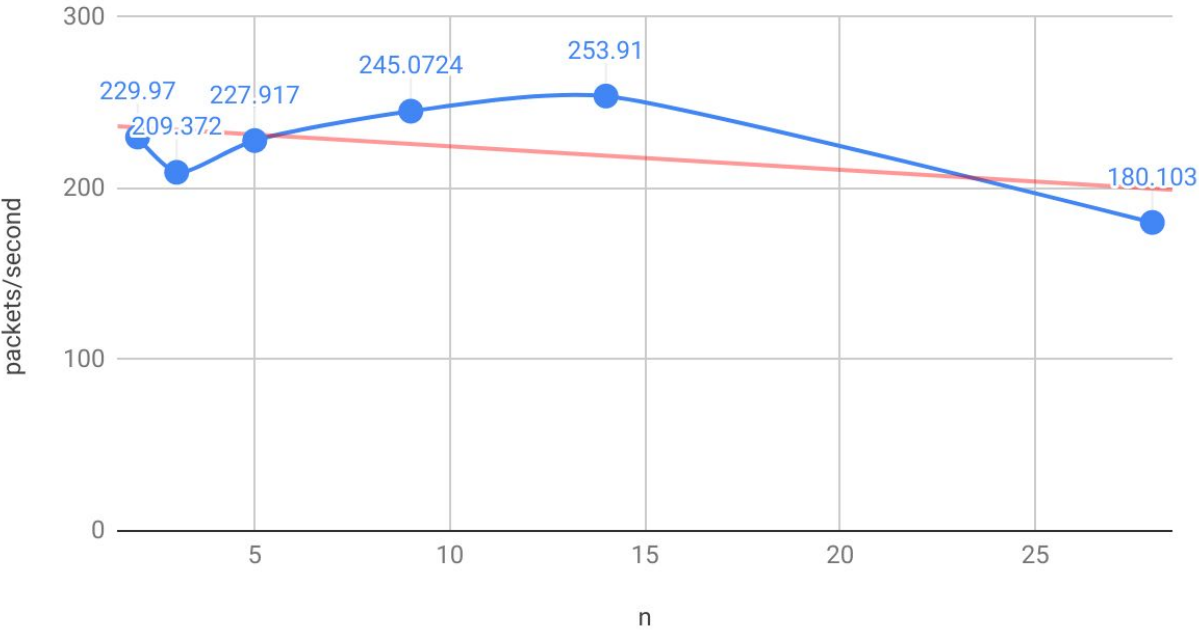
speedup exponential packets with  $w = 1000$  and  $n = 2$  is 3.399936  
speedup exponential packets with  $w = 1000$  and  $n = 3$  is 3.703270  
speedup exponential packets with  $w = 1000$  and  $n = 5$  is 4.296386  
speedup exponential packets with  $w = 1000$  and  $n = 9$  is 4.522426  
speedup exponential packets with  $w = 1000$  and  $n = 14$  is 4.326938  
speedup exponential packets with  $w = 1000$  and  $n = 28$  is 2.478234  
speedup exponential packets with  $w = 2000$  and  $n = 2$  is 6.627845  
speedup exponential packets with  $w = 2000$  and  $n = 3$  is 6.906067  
speedup exponential packets with  $w = 2000$  and  $n = 5$  is 8.147613  
speedup exponential packets with  $w = 2000$  and  $n = 9$  is 8.275903  
speedup exponential packets with  $w = 2000$  and  $n = 14$  is 8.624423  
speedup exponential packets with  $w = 2000$  and  $n = 28$  is 6.190574  
speedup exponential packets with  $w = 4000$  and  $n = 2$  is 13.983295  
speedup exponential packets with  $w = 4000$  and  $n = 3$  is 13.482806  
speedup exponential packets with  $w = 4000$  and  $n = 5$  is 14.692791  
speedup exponential packets with  $w = 4000$  and  $n = 9$  is 16.763318  
speedup exponential packets with  $w = 4000$  and  $n = 14$  is 16.805151  
speedup exponential packets with  $w = 4000$  and  $n = 28$  is 10.432623  
speedup exponential packets with  $w = 8000$  and  $n = 2$  is 25.911302  
speedup exponential packets with  $w = 8000$  and  $n = 3$  is 27.579255  
speedup exponential packets with  $w = 8000$  and  $n = 5$  is 29.516639  
speedup exponential packets with  $w = 8000$  and  $n = 9$  is 32.175600  
speedup exponential packets with  $w = 8000$  and  $n = 14$  is 33.636067  
speedup exponential packets with  $w = 8000$  and  $n = 28$  is 27.26255



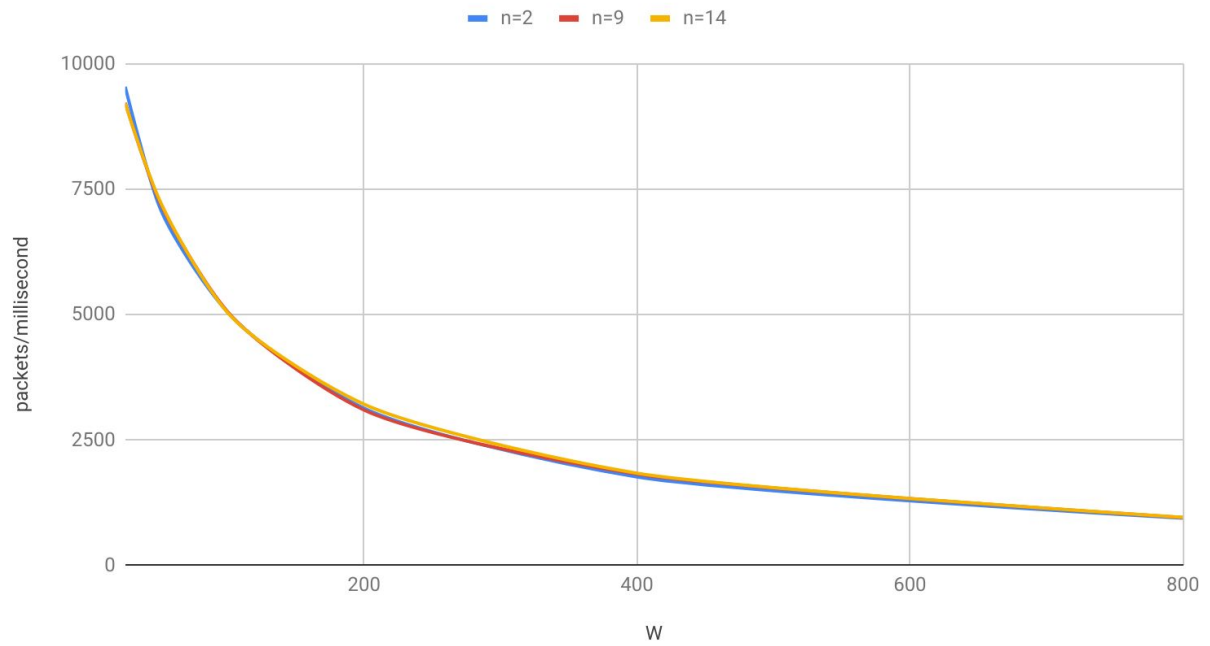
Queue Overhead Speedup



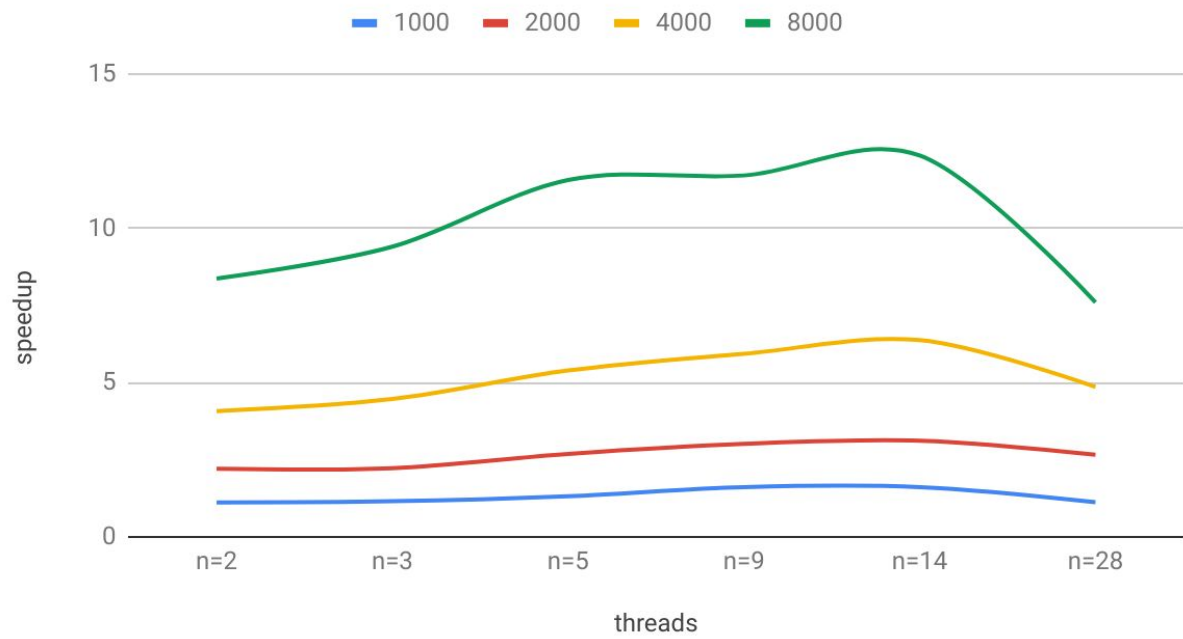
Dispatcher Rate



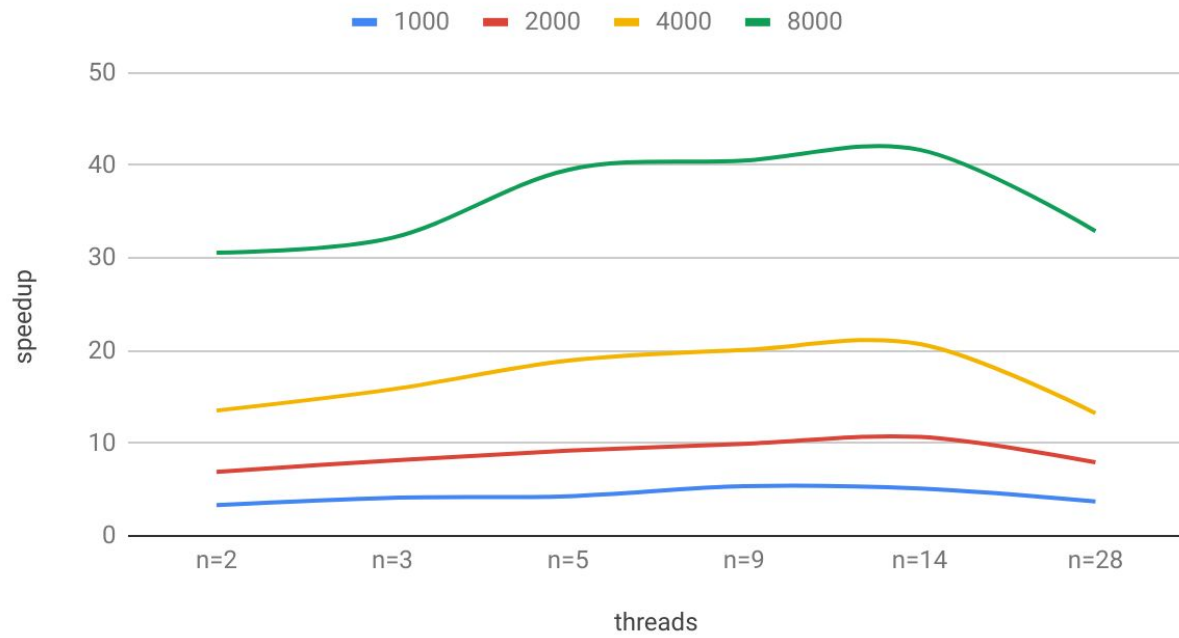
## Worker Rate



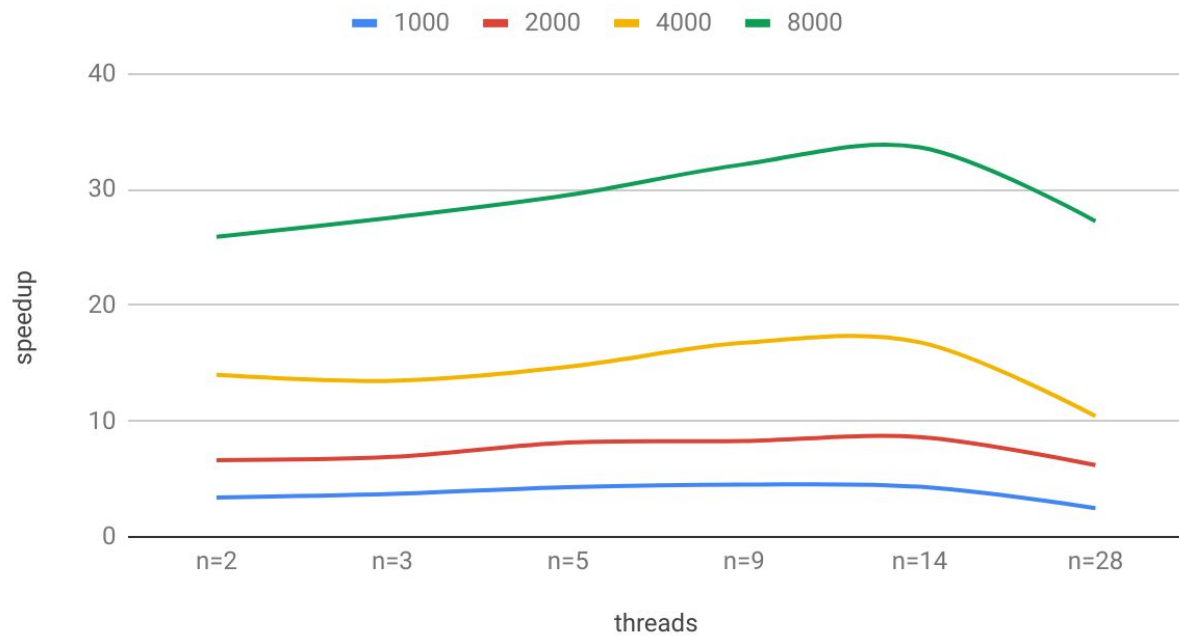
## Constant Packet Speedup



## Uniform Packet Speedup



## Exponential Packet Speedup





- Performance Results

Overall, the results of my performance tests somewhat fall in line with my hypothesis. Starting with queue overhead, I predicted that the extra run time due to the queue would be most pronounced at higher thread counts. However, the overhead due to enqueueing and dequeuing did not change much as more threads were added (aside from some odd results in the beginning).

The dispatcher rate test fell in line with my hypothesis. I predicted that at higher thread counts, the extra overhead to allocate threads and enqueue/dequeue items would cause the dispatcher rate to decrease. The tests saw some immediate gains when the thread count was increased slightly, but when the thread count was 28, the dispatcher rate dropped significantly. This leads me to believe that the queue overhead, as well as the extra work required of the workers and dispatcher to handle the larger number of threads, leads to a significant decrease in performance at higher thread counts.

The worker rate graph shows exactly what I expected. As  $W$  increases, the packets per second goes down significantly. This decrease is not as pronounced in the parallel version however.

I predicted that the speedup between parallel and serial for all the different types of packets would be more or less linear. Also, I suspected that the slope of this linear line would decrease from constant to uniform to exponential. While the graphs showed a very slight linear trend, it was nowhere near what I expected. Interestingly, the speedup dropped significantly at 28 threads. As observed in the dispatcher rate test, very high thread counts introduce a lot of extra overhead. It seems that as thread counts increase, the extra overhead outweighs much of the benefit of the extra parallelization, especially at very high thread counts.

Another interesting result of my test was the increased speedup as the mean work increased. Speedup appears to double for each doubling of the mean work. This is obviously due to this extra work being split among the threads in parallel while being added linearly to the serial version. Also, This change was much more pronounced in the uniform and exponential version when compared to the constant version. This goes against my hypothesis as I suspected that the increased load imbalance would decrease the performance gains as the parallelization would be less effective. I believe this is because the constant packets function I created returned packets that were much less time consuming than the ones returned by the uniform and exponential functions. This would mean that performance gains would be less noticeable as both the parallel and serial versions would finish relatively quickly. Regardless, the exponential packets saw slightly less performance gains than the uniform packets which falls in line with my hypothesis.